

Intro to R - Answers

This lesson is adapted from the Data Carpentry Ecology Lessons 1 (<https://datacarpentry.org/R-ecology-lesson/01-intro-to-r.html>) and Data Carpentry Ecology Lessons 2 (<https://datacarpentry.org/R-ecology-lesson/02-starting-with-data.html>)

This is an R Markdown (<http://rmarkdown.rstudio.com>) Notebook. When you execute code within the notebook, the results appear beneath the code. e.g.

```
3 + 5
```

```
## [1] 8
```

Execute code chunks by clicking the *Run* button within the chunk or by placing your cursor inside it and pressing **Ctrl + Shift + Enter** or OS X: **Cmd + Shift + Enter**. To execute a specific line within a code chunk, placing your cursor on that line and pressing **Ctrl + Enter** or OS X: **Cmd + Enter**. e.g.

Add a new chunk of code in your language of choice by clicking the *Insert Chunk* button on the toolbar or by pressing **Ctrl + Alt + I** or OS X: **Cmd + Option + I**.

```
3 * 2
```

```
## [1] 6
```

TIP: The R Markdown notebook (e.g. *'Intro_to_R.Rmd'*) can be viewed or edited in either in *'Source'* mode which shows the raw Markdown commands, or in the *'Visual'* mode which shows how the Markdown will look in its formatted form. You can switch between the *'Source'* and *'Visual'* modes at any time using the tabs at the top of the notebook.

TIP: When you save the notebook, a HTML file containing the code and output will be saved alongside it (click the Preview button or press Ctrl + Shift + K or OS X: Cmd + Shift + K to preview the HTML file in the *'Viewer'* tab). The preview shows you a rendered HTML copy of the contents of the editor. Consequently, unlike Knit, Preview does not run any R code chunks. Instead, the output of the chunk when it was last run in the editor is displayed.

1. Objects in R

1.1 Creating objects in R

You can get output from R simply by typing math in the console (having space between operators is optional):

```
12 / 7
```

```
## [1] 1.714286
```

However, to do useful and interesting things, we need to assign *values* to *objects*. To create an object, we need to give it a name followed by the assignment operator `<-`, and the value we want to give it:

```
weight_kg <- 55
```

`<-` is the assignment operator. It assigns values on the right to objects on the left. So, after executing `x <- 3`, the value of `x` is `3`. The arrow can be read as **3 goes into x**.

For historical reasons, you can also use `=` for assignments, but not in every context. *Note* `<-` assigns values to your global environment, i.e. they can be used throughout your scripts.

TIP: In RStudio, typing `Alt + -` will write `<-` in a single keystroke in a PC, while typing `Option + -` does the same in a Mac.

1.2 Naming objects (*extra material*)

- objects can have any name, e.g. `x`, `id`, `new_id`
- they **cannot** start with a number, `2x` is not valid, but `x2` is
- R is case sensitive so `id` is different from `ID`
- there are some reserved names, e.g., `if`, `else`, `for` (see here (<https://stat.ethz.ch/R-manual/R-devel/library/base/html/Reserved.html>) for a complete list)
- It's good to avoid dots (`.`) within names. Many function names in R itself have them and dots also have a special meaning (methods) in R and other programming languages.
- recommended to use nouns for object names, and verbs for function names.

1.3 Objects vs. variables

What are known as `objects` in R are known as `variables` in many other programming languages. Depending on the context, `object` and `variable` can have drastically different meanings. However, in this lesson, the two words are used synonymously. For more information see: <https://cran.r-project.org/doc/manuals/r-release/R-lang.html#Objects> (<https://cran.r-project.org/doc/manuals/r-release/R-lang.html#Objects>)

When assigning a value to an object, R does not print anything. You can force R to print the value by using parentheses or by typing the object name:

```
weight_kg <- 55      # doesn't print anything
(weight_kg <- 55)    # but putting parenthesis around the call prints the value of `weight_kg`
```

```
## [1] 55
```

```
weight_kg           # and so does typing the name of the object
```

```
## [1] 55
```

TIP: In the *Environment* tab you can see the current values for variables in memory.

Now that R has `weight_kg` in memory, we can do arithmetic with it. For instance, we may want to convert this weight into pounds (weight in pounds is 2.2 times the weight in kg):

```
2.2 * weight_kg
```

```
## [1] 121
```

We can change an object's value by assigning it a new one:

```
weight_kg <- 57.5
2.2 * weight_kg
```

```
## [1] 126.5
```

This means that assigning a value to one object does not change the values of other objects. For example, let's store the animal's weight in pounds in a new object, `weight_lb`:

```
weight_lb <- 2.2 * weight_kg
```

and then change `weight_kg` to 100.

```
weight_kg <- 100
```

Question: What do you think is the current content of the object `weight_lb`? 126.5 or 220?

Let's see the answer:

```
weight_lb
```

```
## [1] 126.5
```

1.4 Comments

The comment character in R is `#`, anything to the right of a `#` in a script will be ignored by R. It is useful to leave notes and explanations in your scripts.

TIP: RStudio makes it easy to comment or uncomment a paragraph: after selecting the lines you want to comment, press at the same time on your keyboard `Ctrl + Shift + C`. If you only want to comment out one line, you can put the cursor at any location of that line (i.e. no need to select the whole line), then press `Ctrl + Shift + C`.

2. Functions and their arguments

Functions are “canned scripts” that automate more complicated sets of commands including operations, assignments, etc. Many functions are predefined, or can be made available by importing R **packages**.

A function usually takes one or more inputs called **arguments**. Functions often (but not always) return a **value**.

A typical example would be the function `sqrt()`. The input (the argument) must be a number, and the return value (in fact, the output) is the square root of that number. Executing a function (‘running it’) is called **calling** the function. An example of a function call is:

```
ten <- sqrt(weight_kg)
```

Here, the value of `weight_kg` is given to the `sqrt()` function, the `sqrt()` function calculates the square root, and returns the value which is then assigned to the object `ten`. This function is very simple, because it takes just one argument.

The **return ‘value’** of a function need not be numerical (like that of `sqrt()`), and it also does not need to be a single item: it can be a set of things, or even a dataset. We'll see that when we read data files into R.

Arguments can be anything, not only numbers or filenames, but also other objects. Exactly what each argument means differs per function, and must be looked up in the documentation (see below).

Some functions take arguments which may either be specified by the user, or, if left out, take on a *default* value: these are called *options*. Options are typically used to alter the way the function operates, such as whether it ignores ‘bad values’, or what symbol to use in a plot. However, if you want something specific, you can specify a value of your choice which will be used instead of the default.

Let’s try a function that can take multiple arguments: `round()` .

```
round(3.14159)
```

```
## [1] 3
```

Here, we’ve called `round()` with just one argument, `3.14159` , and it has returned the value `3` . That’s because the default is to round to the nearest whole number. If we want more digits we can see how to do that by getting information about the `round` function. We can use `args(round)` to find what arguments it takes, or look at the help for this function using `?round` .

```
args(round)
```

```
## function (x, digits = 0)
## NULL
```

```
?round
```

TIP: In the *Help* tab you can see the help that you just called

We see that if we want a different number of digits, we can type `digits = 2` or however many we want.

```
round(3.14159, digits = 2)
```

```
## [1] 3.14
```

If you provide the arguments in the exact same order as they are defined you don’t have to name them:

```
round(3.14159, 2)
```

```
## [1] 3.14
```

And if you do name the arguments, you can switch their order:

```
round(digits = 2, x = 3.14159)
```

```
## [1] 3.14
```

TIP: It’s good practice to put the non-optional arguments (like the number you’re rounding) first in your function call, and to then specify the names of all optional arguments. If you don’t, someone reading your code might have to look up the definition of a function with unfamiliar arguments to understand what you’re doing.

3. Vectors and data types

3.1. Defining vectors

A vector is the most common and basic data type in R, and is pretty much the workhorse of R. A vector is composed by a series of values, which can be either numbers or characters. We can assign a series of values to a vector using the `c()` function. For example we can create a vector of animal weights and assign it to a new object `weight_g`:

```
weight_g <- c(50, 60, 65, 82)
weight_g
```

```
## [1] 50 60 65 82
```

A vector can also contain characters:

```
animals <- c("mouse", "rat", "dog")
animals
```

```
## [1] "mouse" "rat"    "dog"
```

The quotes around “mouse”, “rat”, etc. are essential here. Without the quotes R will assume objects have been created called `mouse`, `rat` and `dog`. As these objects don’t exist in R’s memory, there will be an error message.

There are many functions that allow you to inspect the content of a vector. `length()` tells you how many elements are in a particular vector:

```
length(weight_g)
```

```
## [1] 4
```

```
length(animals)
```

```
## [1] 3
```

3.2 Vector types

An important feature of a vector, is that all of the elements are the same type of data. The function `class()` indicates the class (the type of element) of an object:

```
class(weight_g)
```

```
## [1] "numeric"
```

```
class(animals)
```

```
## [1] "character"
```

TIP: In the *Environment* tab you can see the type for vectors in memory.

The function `str()` provides an overview of the structure of an object and its elements. It is a useful function when working with large and complex objects:

```
str(weight_g)
```

```
##  num [1:4] 50 60 65 82
```

```
str(animals)
```

```
##  chr [1:3] "mouse" "rat" "dog"
```

3.3 Adding elements to vectors

You can use the `c()` function to add other elements to your vector:

```
weight_g <- c(weight_g, 90) # add to the end of the vector
weight_g <- c(30, weight_g) # add to the beginning of the vector
weight_g
```

```
## [1] 30 50 60 65 82 90
```

In the first line, we take the original vector `weight_g`, add the value `90` to the end of it, and save the result back into `weight_g`. Then we add the value `30` to the beginning, again saving the result back into `weight_g`.

We can do this over and over again to grow a vector, or assemble a dataset. As we program, this may be useful to add results that we are collecting or calculating.

3.4. Atomic Vectors

An **atomic vector** is the simplest R **data type** and is a linear vector of a single type. Above, we saw 2 of the 6 main **atomic vector** types that R uses: "character" and "numeric" (or "double"). These are the basic building blocks that all R objects are built from. The other 4 **atomic vector** types are:

- "logical" for TRUE and FALSE (the boolean data type)
- "integer" for integer numbers (e.g., `2L`, the `L` indicates to R that it's an integer)
- "complex" to represent complex numbers with real and imaginary parts (e.g., `1 + 4i`) and that's all we're going to say about them
- "raw" for bitstreams that we won't discuss further

You can check the type of your vector using the `typeof()` function and inputting your vector as the argument.

Extra material Vectors are one of the many **data structures** that R uses. Other important ones are lists (`list`), matrices (`matrix`), data frames (`data.frame`), factors (`factor`) and arrays (`array`).

4. Subsetting vectors

4.1. Square bracket notation

If we want to extract one or several values from a vector, we must provide one or several indices in square brackets. For instance:

```
animals <- c("mouse", "rat", "dog", "cat")
animals[2]
```

```
## [1] "rat"
```

```
animals[c(3, 2)]
```

```
## [1] "dog" "rat"
```

Extra material - We can also repeat the indices to create an object with more elements than the original one:

```
more_animals <- animals[c(1, 2, 3, 2, 1, 4)]
more_animals
```

```
## [1] "mouse" "rat" "dog" "rat" "mouse" "cat"
```

4.2. Indexing - *Important!*

R indices start at 1. Programming languages like Fortran, MATLAB, Julia, and R start counting at 1, because that's what human beings typically do.

Languages in the C family (including C++, Java, Perl, and Python) count from 0 because that's simpler for computers to do.

4.3 Conditional subsetting of vectors

Another common way of subsetting is by using a logical vector. `TRUE` will select the element with the same index, while `FALSE` will not:

```
weight_g <- c(21, 34, 39, 54, 55)
weight_g[c(TRUE, FALSE, FALSE, TRUE, TRUE)]
```

```
## [1] 21 54 55
```

Typically, these logical vectors are not typed by hand, but are the output of other functions or logical tests. For instance, if you wanted to select only the values above 50:

```
weight_g > 50    # will return logicals with TRUE for the indices that meet the condition
```

```
## [1] FALSE FALSE FALSE TRUE TRUE
```

```
## so we can use this to select only the values above 50
weight_g[weight_g > 50]
```

```
## [1] 54 55
```

You can combine multiple tests using `&` (both conditions are true, AND) or `|` (at least one of the conditions is true, OR):

```
weight_g[weight_g < 30 | weight_g > 50]
```

```
## [1] 21 54 55
```

Here, `<` stands for “less than”, `>` for “greater than”, `>=` for “greater than or equal to”, and `==` for “equal to”. The double equal sign `==` is a test for numerical equality between the left and right hand sides, and should not be confused with the single `=` sign, which performs variable assignment (similar to `<-`).

Extra material - another example

```
weight_g[weight_g >= 30 & weight_g == 21]
```

```
## numeric(0)
```

4.4 Searching for strings in a vector

A common task is to search for certain strings in a vector. One could use the “or” operator `|` to test for equality to multiple values, but this can quickly become tedious. The function `%in%` allows you to test if any of the elements of a search vector are found:

```
animals <- c("mouse", "rat", "dog", "cat")
animals[animals == "cat" | animals == "rat"] # returns both rat and cat
```

```
## [1] "rat" "cat"
```

```
animals %in% c("rat", "cat", "dog", "duck", "goat")
```

```
## [1] FALSE  TRUE  TRUE  TRUE
```

```
animals[animals %in% c("rat", "cat", "dog", "duck", "goat")]
```

```
## [1] "rat" "dog" "cat"
```

5. Missing data

As R was designed to analyze datasets, it includes the concept of missing data (which is uncommon in other programming languages). Missing data are represented in vectors as `NA`.

When doing operations on numbers, most functions will return `NA` if the data you are working with include missing values. This feature makes it harder to overlook the cases where you are dealing with missing data. You can add the argument `na.rm = TRUE` to calculate the result while ignoring the missing values.

```
heights <- c(2, 4, 4, NA, 6)
max(heights)
```

```
## [1] NA
```

```
max(heights, na.rm = TRUE)
```

```
## [1] 6
```


If your data include missing values, you may want to become familiar with the functions `is.na()`, `na.omit()`, and `complete.cases()`. See below for examples.

```
## Extract those elements which are not missing values.
heights[!is.na(heights)]
```

```
## [1] 2 4 4 6
```

Extra material

```
## Returns the object with incomplete cases removed. The returned object is an atomic
vector of type "numeric" (or "double").
na.omit(heights)
```



```
## [1] 2 4 4 6
## attr("na.action")
## [1] 4
## attr("class")
## [1] "omit"
```

```
## Extract those elements which are complete cases. The returned object is an atomic
vector of type "numeric" (or "double").
heights[complete.cases(heights)]
```

```
## [1] 2 4 4 6
```

Recall that you can use the `typeof()` function to find the type of your atomic vector.

Note that data should only ever be removed if there is a valid reason to do so.

6. Data Frames

6.1 What is a Data Frame?

Data frames are the *de facto* data structure for most tabular data in R, and what we use for statistics and plotting.

A data frame can be created by hand, but most commonly they are generated by the functions `read.csv()` or `read_csv()` (tidyverse version); in other words, when importing spreadsheets from your hard drive (or the web).

A data frame is the representation of data in the format of a table where the columns are vectors that all have the same length. Because columns are vectors, each column must contain a single type of data (e.g., characters, integers, factors). For example, here is a figure depicting a data frame comprising a numeric, a character, and a logical vector.

data frame

1	"S"	TRUE
7	"A"	FALSE
3	"U"	TRUE
numeric	character	logical

6.2 Downloading some data and loading it into a Data Frame

We are going to use the R function `download.file()` to download the CSV file that contains the survey data from Figshare, and we will use `read.csv()` to load into memory the content of the CSV file as an object of class `data.frame`.

Inside the `download.file` command, the first entry is a character string with the source URL (`"https://ndownloader.figshare.com/files/2292169"` (`https://ndownloader.figshare.com/files/2292169`)). This source URL downloads a CSV file from figshare. The text after the comma (`"data_raw/portal_data_joined.csv"`) is the destination of the file on your local machine. You'll need to have a folder on your machine called `"data_raw"` where you'll download the file. So this command downloads a file from Figshare, names it `"portal_data_joined.csv"` and adds it to a preexisting folder named `"data_raw"`.

```
download.file(url = "https://ndownloader.figshare.com/files/2292169",
              destfile = "data_raw/portal_data_joined.csv")
```

You are now ready to load the data into a Data Frame.

```
surveys <- read.csv("data_raw/portal_data_joined.csv", stringsAsFactors = FALSE)
```

This statement doesn't produce any output because, as you might recall, assignments don't display anything. If we want to check that our data has been loaded, we can see the contents of the data frame by typing its name: `surveys`.

surveys

record_id	mo...	d..	y...	plot_id	species_id	s...	hindfoot_length	weight	genus	
<int>	<int>	<int>	<int>	<int>	<chr>	<chr>	<int>	<int>	<chr>	►
1	7	16	1977	2	NL	M	32	NA	Neotoma	
72	8	19	1977	2	NL	M	31	NA	Neotoma	
224	9	13	1977	2	NL		NA	NA	Neotoma	
266	10	16	1977	2	NL		NA	NA	Neotoma	
349	11	12	1977	2	NL		NA	NA	Neotoma	
363	11	12	1977	2	NL		NA	NA	Neotoma	
435	12	10	1977	2	NL		NA	NA	Neotoma	
506	1	8	1978	2	NL		NA	NA	Neotoma	
588	2	18	1978	2	NL	M	NA	218	Neotoma	
661	3	11	1978	2	NL		NA	NA	Neotoma	

1-10 of 10,000 rows | 1-10 of 13 columns

Previous123456...1000Next

Wow... that was a lot of output. At least it means the data loaded properly.

6.3 Inspecting data.frame Objects

Let's check the top (the first 6 lines) of this data frame using the function `head()`:

head(surveys)

	record_id	mo...	d...	year	plot_id	species_id	s...	hindfoot_length	weight	
	<int>	<int>	<int>	<int>	<int>	<chr>	<chr>	<int>	<int>	►
1	1	7	16	1977	2	NL	M	32	NA	
2	72	8	19	1977	2	NL	M	31	NA	
3	224	9	13	1977	2	NL		NA	NA	
4	266	10	16	1977	2	NL		NA	NA	
5	349	11	12	1977	2	NL		NA	NA	
6	363	11	12	1977	2	NL		NA	NA	

6 rows | 1-10 of 14 columns

Try also
View(surveys)

We can see this when inspecting the **structure** of a data frame with the function `str()`:

```
str(surveys)
```

```
## 'data.frame':    34786 obs. of  13 variables:
## $ record_id      : int  1 72 224 266 349 363 435 506 588 661 ...
## $ month          : int  7 8 9 10 11 11 12 1 2 3 ...
## $ day            : int  16 19 13 16 12 12 10 8 18 11 ...
## $ year           : int  1977 1977 1977 1977 1977 1977 1977 1977 1978 1978 1978 ...
## $ plot_id        : int  2 2 2 2 2 2 2 2 2 2 ...
## $ species_id     : chr   "NL" "NL" "NL" "NL" ...
## $ sex            : chr   "M" "M" "" "" ...
## $ hindfoot_length: int   32 31 NA NA NA NA NA NA NA NA ...
## $ weight         : int   NA NA NA NA NA NA NA NA 218 NA ...
## $ genus          : chr   "Neotoma" "Neotoma" "Neotoma" "Neotoma" ...
## $ species        : chr   "albigula" "albigula" "albigula" "albigula" ...
## $ taxa           : chr   "Rodent" "Rodent" "Rodent" "Rodent" ...
## $ plot_type      : chr   "Control" "Control" "Control" "Control" ...
```

6.4 More ways to inspect data.frame Objects (*extra material*)

We already saw how the functions `head()` and `str()` can be useful to check the content and the structure of a data frame. Here is a non-exhaustive list of functions to get a sense of the content/structure of the data. Let's try them out!

Size:

- `dim(surveys)` - returns a vector with the number of rows in the first element, and the number of columns as the second element (the **dimensions** of the object)
- `nrow(surveys)` - returns the number of rows
- `ncol(surveys)` - returns the number of columns

Content:

- `head(surveys)` - shows the first 6 rows
- `tail(surveys)` - shows the last 6 rows

Names:

- `names(surveys)` - returns the column names (synonym of `colnames()` for `data.frame` objects)
- `rownames(surveys)` - returns the row names

Summary:

- `str(surveys)` - structure of the object and information about the class, length and content of each column
- `summary(surveys)` - summary statistics for each column

TIP: Most of these functions are “generic”, they can be used on other types of objects besides `data.frame`.

6.5 Indexing data frames

Our survey data frame has rows and columns (it has 2 dimensions), if we want to extract some specific data from it, we need to specify the “coordinates” we want from it. Row numbers come first, followed by column numbers. However, note that different ways of specifying these coordinates lead to results with different classes.

```
# first element in the first column of the data frame (as a vector)  
surveys[1, 1]
```

```
## [1] 1
```

```
# first element in the 6th column (as a vector)  
surveys[1, 6]
```

```
## [1] "NL"
```

```
# first column of the data frame (as a vector)  
surveys[, 1]
```

##	[1]	1	72	224	266	349	363	435	506	588	661	748	845
##	[13]	990	1164	1261	1374	1453	1756	1818	1882	2133	2184	2406	2728
##	[25]	3000	3002	4667	4859	5048	5180	5299	5485	5558	5583	5966	6020
##	[37]	6023	6036	6167	6479	6500	8022	8263	8387	8394	8407	8514	8543
##	[49]	8657	8675	8755	9048	9132	9270	9349	9512	9605	9703	10606	10617
##	[61]	10627	10720	10923	10949	11215	11329	11496	11498	11611	11628	11709	11877
##	[73]	11879	11887	11953	12299	12458	12602	12678	12708	12729	12756	12871	12997
##	[85]	13025	13114	13275	13326	13434	13465	13476	13615	13630	13748	13977	14646
##	[97]	14740	14796	14804	14896	14984	15302	15780	16407	17230	17877	17941	18110
##	[109]	18364	18486	18639	18678	18680	18846	18948	19032	19245	20040	20089	20364
##	[121]	20466	20992	21113	21120	21161	21232	21235	21286	21507	21558	22314	22728
##	[133]	22899	23032	24589	24690	24703	25247	25701	26028	26644	26827	26857	27042
##	[145]	27326	27734	27781	27860	27897	27919	28022	28077	28328	28551	28556	28662
##	[157]	28928	29039	29310	29435	29572	29823	29865	29950	30066	30067	30172	30174
##	[169]	30175	30307	30308	30309	30443	30741	31516	31517	31633	31717	31721	31862
##	[181]	31946	32045	32167	32391	32817	33040	33041	33237	33238	33339	33415	33583
##	[193]	33586	33837	33847	33966	34198	34783	34991	35212	35404	3	226	233
##	[205]	245	251	257	259	268	346	350	354	361	422	424	427
##	[217]	438	449	507	511	518	519	521	522	523	527	533	593
##	[229]	595	600	601	660	664	733	754	757	764	838	842	849
##	[241]	914	924	930	1018	1061	1062	1159	1183	1269	1272	1277	1370
##	[253]	1382	1461	1469	1519	1529	1534	1535	1539	1581	1611	1644	1661
##	[265]	1676	1692	1701	1768	1839	1845	1846	2014	2060	2070	2122	2131
##	[277]	2176	2181	2338	2339	2340	2372	2400	2456	2460	2525	2555	2589
##	[289]	2600	2622	2631	2635	2643	2653	2662	2667	2692	2695	2701	2715
##	[301]	2740	2790	2814	2901	2902	2906	2975	2981	3033	3035	3099	3107
##	[313]	3182	3236	3302	3388	3454	3456	3525	3547	3550	3564	3566	3569
##	[325]	3580	3738	3741	3781	3790	3867	3877	3889	3897	3899	3901	3947
##	[337]	3960	3979	3998	4004	4198	4238	4243	4251	4268	4271	4276	4306
##	[349]	4317	4320	4341	4347	4452	4482	4520	4548	4685	4693	4714	4952
##	[361]	5044	5063	5167	5288	5483	6006	6184	6470	6474	6491	6520	6523
##	[373]	6528	6549	6553	6824	6832	6837	6842	6889	6939	6944	6945	7061
##	[385]	7080	7246	7294	7402	7420	7553	7558	7715	7723	7735	7884	7895
##	[397]	7916	8125	8235	8339	8359	8413	8494	8631	8656	8800	8903	8954
##	[409]	8961	9034	9133	9156	9180	9182	9269	9452	9572	9672	9762	9767
##	[421]	9873	9880	9894	9901	9920	9980	9984	10001	10004	10006	10021	10031
##	[433]	10092	10109	10114	10254	10277	10290	10318	10354	10363	10366	10411	10415
##	[445]	10421	10430	10447	10500	10517	10518	10532	10595	10600	10603	10610	10621
##	[457]	10623	10686	10693	10716	10784	10802	10804	10899	10909	10920	11013	11015
##	[469]	11063	11074	11092	11148	11150	11161	11186	11187	11197	11210	11284	11291
##	[481]	11304	11305	11316	11327	11387	11388	11391	11408	11421	11489	11514	11515
##	[493]	11535	11584	11595	11600	11608	11684	11695	11722	11772	11784	11788	11791
##	[505]	11841	11862	11920	12011	12014	12056	12117	12137	12138	12234	12249	12389
##	[517]	12394	12407	12419	12439	12533	12535	12545	12676	12764	12780	12966	12972
##	[529]	13112	13145	13167	13267	13481	13589	13757	13820	13822	13920	13922	13935
##	[541]	14098	14104	14119	14123	14134	14258	14279	14322	14423	14491	14540	14603
##	[553]	14697	14787	14794	14797	14889	14988	16292	16377	17298	17320	17331	17910
##	[565]	17915	17973	18096	18193	18201	18654	18743	18821	19378	19427	19539	19806
##	[577]	19866	20856	21167	21215	21223	21269	21275	21284	21325	21385	21435	21442
##	[589]	21448	21502	21505	21555	21560	21708	21757	21814	21879	21940	22003	22042
##	[601]	22044	22105	22109	22168	22250	22368	22375	22444	22550	22560	22689	22842
##	[613]	22853	22865	22880	22896	22997	23002	23003	23041	23044	23115	23117	23120
##	[625]	23136	23144	23156	23169	23220	23225	23234	23237	23243	23257	23341	23345
##	[637]	23357	23365	23368	23403	23498	23519	23522	23549	23565	23645	23650	23687
##	[649]	23703	23895	23899	23902	23904	23922	23929	23936	24027	24045	24197	24292

```
## [34261] 22460 22570 22733 23506 23543 23553 23556 23568 23577 23697 23910 23941
## [34273] 23943 24016 24035 24042 24060 24065 24081 24184 24188 24201 24204 24207
## [34285] 24214 24310 24439 24578 25497 26040 26089 26110 26119 26332 26402 26414
## [34297] 26448 26474 26523 26530 26598 26606 26621 26628 26629 26631 26635 26852
## [34309] 27726 27787 27800 27904 27915 28025 28057 28195 28331 28334 29488 29764
## [34321] 29765 29766 31157 31268 32226 32473 32474 32475 32476 32686 32687 32688
## [34333] 33104 33105 33308 34045 34047 34282 34284 34285 34286 34521 34522 34523
## [34345] 34525 34696 34861 34862 2196 23163 24944 24996 30356 31663 31664 31747
## [34357] 31988 33306 33307 34283 34520 2471 31748 33567 33568 33569 33570 35548
## [34369] 10427 5603 5801 5984 5993 6882 7390 7459 7722 7775 7778 7907
## [34381] 7926 7927 8150 8227 8368 8377 8665 8864 9339 9434 9587 9709
## [34393] 9924 10019 10133 10267 10439 10448 10494 10501 10511 11159 11167 11339
## [34405] 11390 11413 11495 11537 11599 11619 11698 11723 12153 12238 12262 12392
## [34417] 12401 12431 12451 12538 12609 12760 12835 12885 12989 12999 13029 13140
## [34429] 13177 13342 13580 13635 13804 13813 13942 13994 14102 14105 14111 14314
## [34441] 14381 14390 14523 14617 14627 15113 15123 15240 15253 15288 15396 15398
## [34453] 15400 15408 15455 15596 15656 15665 15673 15682 15695 15716 15932 15997
## [34465] 16157 16185 16216 16223 16229 16236 16567 16583 16594 16621 16636 16733
## [34477] 16737 16752 16773 16784 16794 16902 16919 17001 17030 17035 17119 17129
## [34489] 17139 17250 17285 17389 17391 17404 17490 17501 17581 17600 17607 17820
## [34501] 17887 17888 17889 17903 17971 17988 17995 18005 18012 18079 18095 18100
## [34513] 18102 18111 18190 18206 18233 18330 18337 18353 18358 18362 18373 18455
## [34525] 18603 18609 18638 18651 18653 18668 18670 18681 18736 18754 18763 18769
## [34537] 18829 18840 18844 18928 19118 19200 19219 19226 19235 19243 19375 19383
## [34549] 19385 19387 19389 19417 19435 19458 19556 19582 19637 19675 19772 19774
## [34561] 19795 19815 19874 19875 19889 20018 20043 20085 20112 20371 20430 20432
## [34573] 20465 20542 20591 20605 20648 20661 20693 20698 20727 20790 21000 21034
## [34585] 21109 21157 21231 21238 21289 21398 21564 21839 21876 22087 22124 22181
## [34597] 22239 22249 22364 22384 22448 22451 22549 22701 23045 23256 23366 23540
## [34609] 23587 23601 23657 23709 23725 24941 25263 25420 25432 25521 25738 25824
## [34621] 26044 26135 26826 27030 27174 27284 27314 27328 2193 5596 5763 5780
## [34633] 5792 6199 6426 8246 8393 8643 8644 8648 8781 8909 12412 13111
## [34645] 13183 13278 13311 13426 13758 13964 14168 14262 15129 23394 23564 23646
## [34657] 23649 23901 24544 24718 24869 24986 24997 25171 25749 25796 26051 26058
## [34669] 26081 26130 26299 26348 26444 26602 26837 26859 26867 26963 26967 26972
## [34681] 26976 26982 27015 27262 27273 27308 27346 27349 27419 27594 28586 28587
## [34693] 28588 28589 28590 28668 28805 29489 29898 11849 12766 12874 15588 16230
## [34705] 16242 17828 17836 18761 18929 20104 20994 21163 21330 21758 23714 26056
## [34717] 27144 27203 28333 28335 29623 29992 30218 32086 34287 5267 9248 11200
## [34729] 11222 12293 12418 12444 13442 13775 13790 15432 15475 16566 16769 16777
## [34741] 19413 19669 19891 20026 20435 20533 21434 22111 13458 13471 20183 28970
## [34753] 15946 16652 19632 32227 21209 25710 26042 26096 26356 26475 26546 26776
## [34765] 26819 28332 28336 28337 28338 28585 28667 29231 30355 32085 32477 33103
## [34777] 33305 34524 35382 26557 26787 26966 27185 27792 28806 30986
```

```
# first column of the data frame (as a data.frame)
surveys[1]
```

record_id
<int>

1

72

	record_id
	<int>
	224
	266
	349
	363
	435
	506
	588
	661
1-10 of 10,000 rows	Previous 1 2 3 4 5 6 ... 1000 Next

```
# the 3rd row of the data frame (as a data.frame)
surveys[3, ]
```

	record_id	mo...	d...	year	plot_id	species_id	s...	hindfoot_length	weight
	<int>	<int>	<int>	<int>	<int>	<chr>	<chr>	<int>	<int>
3	224	9	13	1977	2	NL		NA	NA
1 row 1-10 of 14 columns									

: is a special function that creates numeric vectors of integers in increasing or decreasing order.

```
# first three elements in the 7th column (as a vector)
surveys[1:3, 7]
```

```
## [1] "M" "M" ""
```

Extra material - Test 1:10 and 10:1 for instance. You can also exclude certain indices of a data frame using the “ - ” sign:

```
surveys[, -1]           # The whole data frame, except the first column
```

mo...	d..	y...	plot_id	species_id	s...	hindfoot_length	weight	genus	species
<int>	<int>	<int>	<int>	<chr>	<chr>	<int>	<int>	<chr>	<chr>
7	16	1977	2	NL	M	32	NA	Neotoma	albigula
8	19	1977	2	NL	M	31	NA	Neotoma	albigula
9	13	1977	2	NL		NA	NA	Neotoma	albigula
10	16	1977	2	NL		NA	NA	Neotoma	albigula
11	12	1977	2	NL		NA	NA	Neotoma	albigula
11	12	1977	2	NL		NA	NA	Neotoma	albigula

mo...	d..	y...	plot_id	species_id	s...	hindfoot_length	weight	genus	species					
<int>	<int>	<int>	<int>	<chr>	<chr>	<int>	<int>	<chr>	<chr>					
12	10	1977	2	NL		NA	NA	Neotoma	albigula					
1	8	1978	2	NL		NA	NA	Neotoma	albigula					
2	18	1978	2	NL	M	NA	218	Neotoma	albigula					
3	11	1978	2	NL		NA	NA	Neotoma	albigula					
1-10 of 10,000 rows 1-10 of 12 columns						Previous	1	2	3	4	5	6	... 1000	Next

```
surveys[-c(7:34786), ] # Equivalent to head(surveys)
```

	record_id	mo...	d...	year	plot_id	species_id	s...	hindfoot_length	weight	
	<int>	<int>	<int>	<int>	<int>	<chr>	<chr>	<int>	<int>	►
1	1	7	16	1977	2	NL	M	32	NA	
2	72	8	19	1977	2	NL	M	31	NA	
3	224	9	13	1977	2	NL		NA	NA	
4	266	10	16	1977	2	NL		NA	NA	
5	349	11	12	1977	2	NL		NA	NA	
6	363	11	12	1977	2	NL		NA	NA	
6 rows 1-10 of 14 columns										

6.6 Subsetting data frames

Data frames can be subset by calling indices (as shown previously), but also by calling their column names directly:

```
surveys["species_id"]      # Result is a data.frame
surveys[, "species_id"]    # Result is a vector
surveys[["species_id"]]    # Result is a vector
surveys$species_id         # Result is a vector
```

TIP: In RStudio, you can use the autocompletion feature to get the full and correct names of the columns.

Extra material Like we've seen in the section on vectors, data frames can also be subset using logical tests:

```
females <- surveys[surveys$sex=="F",]
head(females)
```

	record_id	mo...	d...	year	plot_id	species_id	s...	hindfoot_length	weight	
	<int>	<int>	<int>	<int>	<int>	<chr>	<chr>	<int>	<int>	►
21	2133	10	25	1979	2	NL	F	33	274	
22	2184	11	17	1979	2	NL	F	30	186	
23	2406	1	16	1980	2	NL	F	33	184	

	record_id <int>	mo... <int>	d... <int>	year <int>	plot_id <int>	species_id <chr>	s... <chr>	hindfoot_length <int>	weight <int>
24	2728	3	9	1980	2	NL	F	NA	NA
25	3000	5	18	1980	2	NL	F	31	87
26	3002	5	18	1980	2	NL	F	33	174

6 rows | 1-10 of 14 columns

7. Challenges - Now it's your turn!

7.1 Assigning values to objects

What are the values after each statement in the following?

```
mass <- 47.5           # mass?
# mass is 47.5
age <- 122             # age?
# age is 122
mass <- mass * 2.0     # mass?
# mass is 95
age <- age - 20        # age?
# age is 102
mass_index <- mass/age # mass_index?
# mass_index is 95/102 = 0.93137
```

7.2 Vectors

We've seen that atomic vectors can be of type character, numeric (or double), integer, and logical. But what happens if we try to mix these types in a single vector?

R implicitly converts them to all be the same type

What will happen in each of these examples? (hint: use `class()` to check the data type of your objects):

```
num_char <- c(1, 2, 3, "a")
# coerced to type chr
num_logical <- c(1, 2, 3, TRUE)
# coerced to type num
char_logical <- c("a", "b", "c", TRUE)
# coerced to type chr
tricky <- c(1, 2, 3, "4")
# coerced to type chr
```

Why do you think it happens?

Vectors can be of only one data type. R tries to convert (coerce) the content of this vector to find a “common denominator” that doesn't lose any information.

How many values in `combined_logical` are "TRUE" (as a character) in the following example (reusing the 2 `..._logical`s from above):

```
combined_logical <- c(num_logical, char_logical)
```

Only one. There is no memory of past data types, and the coercion happens the first time the vector is evaluated. Therefore, the `TRUE` in `num_logical` gets converted into a `1` before it gets converted into `"1"` in `combined_logical`.

In R, we call converting objects from one class into another class *coercion*. These conversions happen according to a hierarchy, whereby some types get preferentially coerced into other types. From the examples above, can you draw a diagram that represents the hierarchy of how these data types are coerced?

```
logical → numeric → character ← logical
```

7.3 Logical test

```
# Reminder of logical tests:  
4>5 # FALSE
```

```
## [1] FALSE
```

```
"i" == "I" # FALSE
```

```
## [1] FALSE
```

```
5 == "five" # FALSE
```

```
## [1] FALSE
```

```
"i" %in% letters # TRUE - letters is a vector of the alphabet in lower case
```

```
## [1] TRUE
```

Can you figure out why `"four" > "five"` returns `TRUE`?

```
"four" > "five"
```

```
## [1] TRUE
```

When using `>` or `<` on strings, R compares their alphabetical order. Here `"four"` comes after `"five"`, and therefore is `"greater than"` it.

7.4 Subsetting/Removing NAs

1. Using this vector of heights in inches, create a new vector, `heights_no_na`, with the NAs removed.

```
heights <- c(63, 69, 60, 65, NA, 68, 61, 70, 61, 59, 64, 69, 63, 63, NA, 72, 65, 64, 70, 63, 65)
```

2. Use the function `median()` to calculate the median of the `heights` vector.

3. Use R to figure out how many people in the set are taller than 67 inches.

```
heights <- c(63, 69, 60, 65, NA, 68, 61, 70, 61, 59, 64, 69, 63, 63, NA, 72, 65, 64, 70, 63, 65)
```

```
# 1.
heights_no_na <- heights[!is.na(heights)]
# or
heights_no_na <- na.omit(heights)
# or
heights_no_na <- heights[complete.cases(heights)]

# 2.
median(heights, na.rm = TRUE) # 64
```

```
## [1] 64
```

```
# 3.
heights_above_67 <- heights_no_na[heights_no_na > 67]
length(heights_above_67) # 6
```

```
## [1] 6
```

7.5 Data.frames

Based on the output of `str(surveys)`, can you answer the following questions?

- What is the class of the object `surveys` ?
- How many rows and how many columns are in this object?
- How many species have been recorded during these surveys?

```
str(surveys)
```

```
## 'data.frame':    34786 obs. of  13 variables:
## $ record_id      : int  1 72 224 266 349 363 435 506 588 661 ...
## $ month          : int  7 8 9 10 11 11 12 1 2 3 ...
## $ day            : int  16 19 13 16 12 12 10 8 18 11 ...
## $ year           : int  1977 1977 1977 1977 1977 1977 1977 1977 1978 1978 1978 ...
## $ plot_id        : int  2 2 2 2 2 2 2 2 2 2 ...
## $ species_id     : chr  "NL" "NL" "NL" "NL" ...
## $ sex            : chr  "M" "M" "" "" ...
## $ hindfoot_length: int  32 31 NA NA NA NA NA NA NA NA ...
## $ weight         : int  NA NA NA NA NA NA NA NA 218 NA ...
## $ genus          : chr  "Neotoma" "Neotoma" "Neotoma" "Neotoma" ...
## $ species        : chr  "albigula" "albigula" "albigula" "albigula" ...
## $ taxa           : chr  "Rodent" "Rodent" "Rodent" "Rodent" ...
## $ plot_type      : chr  "Control" "Control" "Control" "Control" ...
```

```
## * class: data frame
## * how many rows: 34786,  how many columns: 13
## * how many species: 48
```

Subsetting

1. Create a `data.frame` (`surveys_200`) containing only the data in row 200 of the `surveys` dataset.
2. Notice how `nrow()` gave you the number of rows in a `data.frame` ?
 - Use that number to pull out just that last row in the data frame.
 - Compare that with what you see as the last row using `tail()` to make sure it's meeting expectations.
 - Pull out that last row using `nrow()` instead of the row number.
 - Create a new data frame (`surveys_last`) from that last row.
3. Use `nrow()` to extract the row that is in the middle of the data frame. Store the content of this row in an object named `surveys_middle`.
4. Combine `nrow()` with the `-` notation above to reproduce the behavior of `head(surveys)`, keeping just the first through 6th rows of the `surveys` dataset.

```
## 1.
surveys_200 <- surveys[200, ]
## 2.
# Saving `n_rows` to improve readability and reduce duplication
n_rows <- nrow(surveys)
surveys_last <- surveys[n_rows, ]
## 3.
surveys_middle <- surveys[n_rows / 2, ]
## 4.
surveys_head <- surveys[-(7:n_rows), ]
```